

Ethan Zaruba-Walker

ethan@staffroomai.com | (512) 694-4844 | GitHub

September 16, 2026

Dear H Company hiring team,

I am writing about the Founding Forward Deployed Engineer role in the US. I live in Emeryville, so the Bay Area seat is already home, and I can travel for on-site deployments. For three years I have been the technical person on customer-facing AI work: name the multi-step process worth automating, integrate it with the systems the team already runs, and stay until they can operate it.

H is an agentic-model lab. That is why I am applying. You train computer-use models so agents can act on a screen when there is no API, and Forward Deployed is how those agents land in a customer's environment and how field findings get back into the H platform. I already run that loop: discover, integrate, iterate, hand over.

How that maps to what you need:

- **Customer-facing AI, three years:** I source, scope, and ship every engagement, including an MSP pre-sales risk workflow cut 98% in processing time, then train operators to extend it.
- **Python production systems:** 7+ years of professional Python, plus REST/OAuth APIs, PostgreSQL, Docker, and data pipelines (ingestion throughput up 99.8% at 86% lower cost). TypeScript and Next.js (React) are in production on a public OAuth app.
- **LLM and agentic workflows:** Claude, OpenAI, and Gemini tool-calling, MCP servers, and Claude Code as a daily driver, orchestrated in n8n when the operating team needed a tool they could keep.

One real gap, named plainly: I have not shipped computer-use agents that click through a GUI. My production agents call APIs, databases, and MCP tools. FastAPI is not my daily framework; Flask and Next.js backends are. The Holo action layer is the part I would ramp on with your research and product teams.

Thank you for your time and consideration. I would welcome a conversation about which US customer workflows a founding FDE would own first.

Sincerely,

Ethan Zaruba-Walker